

EXP NO: 3	Implement of Matrix Multiplication with Hadoop Map Reduce
DATE:	

AIM: -

BACKGROUND THEORY: -

PROCEDURE: -

- Switch to superuser mode using `sudo su`.
- Input Format:
 - We'll have two input files representing matrices A and B.
 - Matrix A (m x n): Split by rows.
 - Matrix B (n x p): Split by columns.
- Map Step:
 - For each element in matrix A (i, k, A[i][k]), emit the product to the intermediate key (i, j) where j is the column index in matrix B.
 - For each element in matrix B (k, j, B[k][j]), emit the product to the intermediate key (i, j) for all rows i of matrix A.
- Reduce Step:
 - For each intermediate key (i, j), sum the products of the corresponding values to calculate the result matrix C (i, j).
- Output Format:
 - Output matrix C, where each line is of the form i, j, C[i][j].

CODING: -

- `sudo su`
- `MatrixMapper.java`
 - `import java.io.IOException;`
 - `import org.apache.hadoop.io.IntWritable;`
 - `import org.apache.hadoop.io.Text;`
 - `import org.apache.hadoop.mapreduce.Mapper;`
 -
 - `public class MatrixMapper extends Mapper<Object, Text, Text, Text> {`
 -
 - `@Override`
 - `public void map(Object key, Text value, Context context) throws IOException, InterruptedException {`
 - `String[] line = value.toString().split(",");`
 - `String matrixName = line[0]; // A or B`
 - `int i = Integer.parseInt(line[1]);`
 - `int j = Integer.parseInt(line[2]);`
 - `int valueOfElement = Integer.parseInt(line[3]);`

-
- if (matrixName.equals("A")) {
- // Emit for all columns of B
- for (int k = 0; k < context.getConfiguration().getInt("p", 0); k++) {
- context.write(new Text(i + "," + k), new Text("A," + j + "," + valueOfElement));
- }
- } else {
- // Emit for all rows of A
- for (int k = 0; k < context.getConfiguration().getInt("m", 0); k++) {
- context.write(new Text(k + "," + j), new Text("B," + i + "," + valueOfElement));
- }
- }
- }
- }
- }

- MatrixReducer.java

- import java.io.IOException;
- import java.util.HashMap;
- import java.util.Map;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Reducer;
-
- public class MatrixReducer extends Reducer<Text, Text, Text, IntWritable> {
-
- @Override
- public void reduce(Text key, Iterable<Text> values, Context context) throws IOException, InterruptedException {
- Map<Integer, Integer> mapA = new HashMap<>();

- Map<Integer, Integer> mapB = new HashMap<>();
-
- for (Text val : values) {
- String[] parts = val.toString().split(",");
- if (parts[0].equals("A")) {
- mapA.put(Integer.parseInt(parts[1]), Integer.parseInt(parts[2]));
- } else {
- mapB.put(Integer.parseInt(parts[1]), Integer.parseInt(parts[2]));
- }
- }
-
- int result = 0;
- for (Integer k : mapA.keySet()) {
- if (mapB.containsKey(k)) {
- result += mapA.get(k) * mapB.get(k);
- }
- }
-
- context.write(key, new IntWritable(result));
- }
- }

- MatrixMultiplication.java

- import org.apache.hadoop.conf.Configuration;
- import org.apache.hadoop.fs.Path;
- import org.apache.hadoop.io.IntWritable;
- import org.apache.hadoop.io.Text;
- import org.apache.hadoop.mapreduce.Job;
- import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
- import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

```

○ public class MatrixMultiplication {
○
○     public static void main(String[] args) throws Exception {
○         Configuration conf = new Configuration();
○         // Dimensions of the matrices A (m x n) and B (n x p)
○         conf.setInt("m", 3); // Rows of A
○         conf.setInt("n", 2); // Columns of A and Rows of B
○         conf.setInt("p", 3); // Columns of B
○
○
○         Job job = Job.getInstance(conf, "Matrix Multiplication");
○         job.setJarByClass(MatrixMultiplication.class);
○         job.setMapperClass(MatrixMapper.class);
○         job.setReducerClass(MatrixReducer.class);
○
○
○         job.setOutputKeyClass(Text.class);
○         job.setOutputValueClass(IntWritable.class);
○
○
○         FileInputFormat.addInputPath(job, new Path(args[0]));
○         FileOutputFormat.setOutputPath(job, new Path(args[1]));
○         System.exit(job.waitForCompletion(true) ? 0 : 1);
○     }
○ }

```

- Steps to Run the Code

```

○ hdfs dfs -put matrixA.txt /input/
○ hdfs dfs -put matrixB.txt /input/
○ hadoop com.sun.tools.javac.Main MatrixMultiplication.java
○ jar cf matrixmultiplication.jar MatrixMultiplication*.class
○ hadoop jar matrixmultiplication.jar MatrixMultiplication /input/ /output/
○ hdfs dfs -cat /output/part-r-00000

```

OUTPUT: -

```
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.4.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.3.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Null.java
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.8.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.5.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.8.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs_0.22.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.9.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs_0.20.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0-alpha4.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.2.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.9.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.0-alpha2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.2.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_2.10.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.0.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.1.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.2.4.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs_0.21.0.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/Apache_Hadoop_HDFS_3.1.3.xml
hadoop-3.3.6/share/hadoop/hdfs/jdiff/hadoop-hdfs-client-3.3.6-tests.jar
hadoop-3.3.6/share/hadoop/hdfs/hadoop-hdfs-https-3.3.6.jar
[root@ip-172-31-40-166 local]# sudo mv hadoop-3.3.6 hadoop
[root@ip-172-31-40-166 local]# sudo nano ~/.bashrc
[root@ip-172-31-40-166 local]# sudo nano ~/.bashrc
[root@ip-172-31-40-166 local]# sudo nano ~/.bashrc
[root@ip-172-31-40-166 local]# source ~/.bashrc
[root@ip-172-31-40-166 local]# hadoop version
hadoop 3.3.6
Source code repository https://github.com/apache/hadoop.git -r 1be78238728da9266a4f8b195058f08d012bf9c
Compiled by ubuntu on 2023-06-18T08:22X
Compiled on platform linux-x86_64
Compiled with protoc 3.7.1
From source with checksum 5652179ad5576cb287d9dc633bb53bbd
This command was run using /usr/local/hadoop/share/hadoop/common/hadoop-common-3.3.6.jar
[root@ip-172-31-40-166 local]#
```

[illegible]